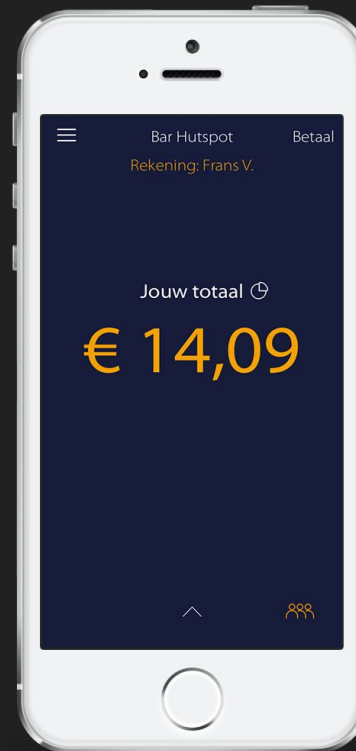
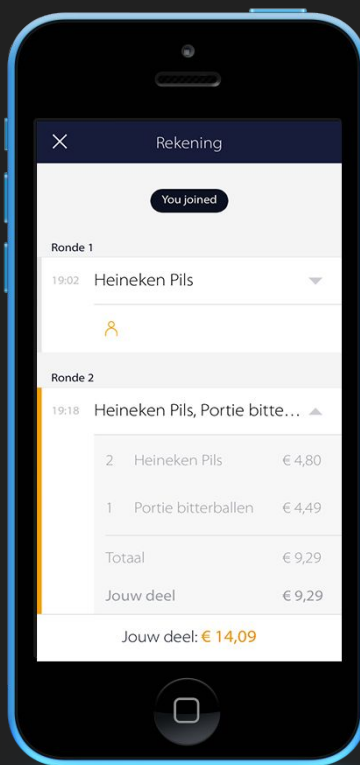


Distributed queues with RabbitMQ

IN4391 Demo Day - April 6, 2016

Jorik Oostenbrink
Dominik Harz


Tabster



Event based architecture




Denormalized read store
Subscribes to events
on the Write side



Query


Read side



Query response




User views data
in the UI



User makes a change
in the UI

Command

Write side




Append
events



Event store
Publishes events after
they have been saved

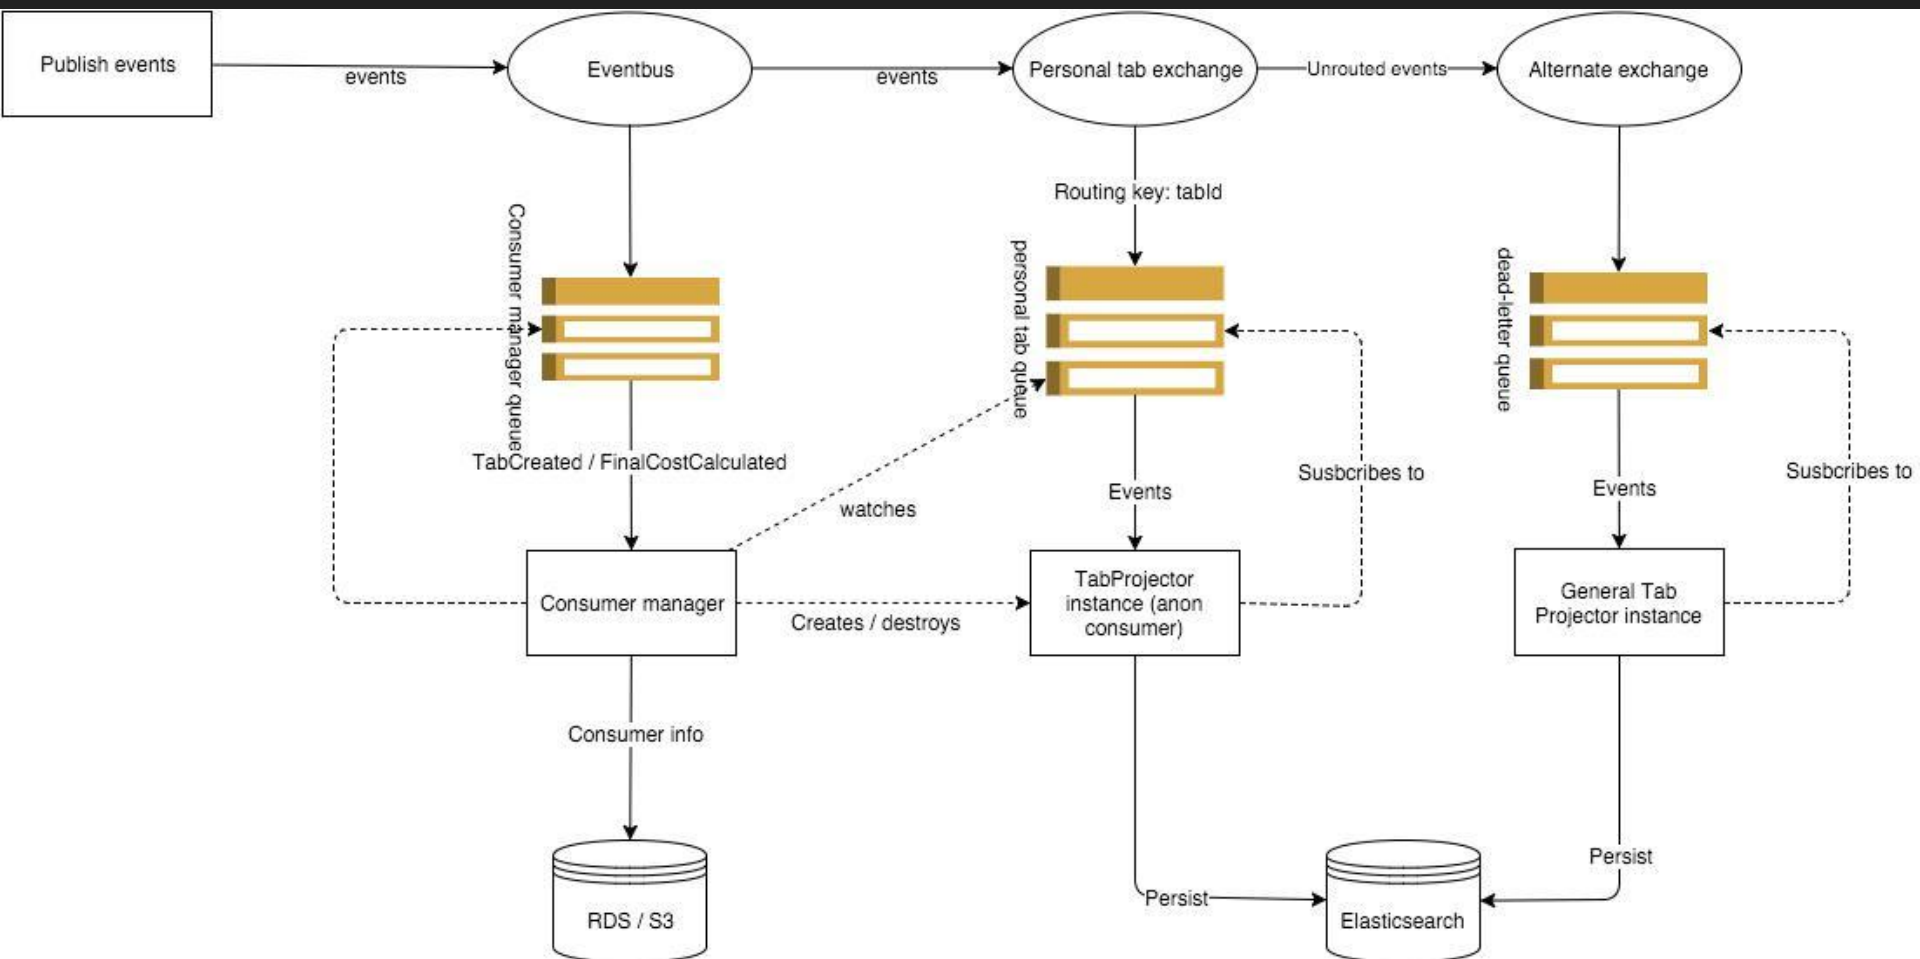
Events

Event based architecture



Key requirements

- A failure of one queue or worker does not affect other system components
- All system events must be logged in the order they occur
- The properties of the Tabster system must be demonstrated when it contains
 - 10 tabs with at least 100 events per tab
 - 2 consumer managers
 - workload consists of at least 1000 events in total



Features

- Eventual consistency
 - Separate read and write models
 - Processes events from separate tabs simultaneously
 - Processes events from a single tab in the order they arrived
- Fault tolerance
 - Queues get recreated after failing (set to durable)
 - Workers are monitored for failures
 - Elasticsearch configured as a cluster with replication
 - API configured as a HA cluster with nginx as load-balancer

Demo time!

Next steps

1. Execute RabbitMQ as a cluster with multiple nodes
2. Place the system on AWS
3. Execute stress testing
 - a. Optimize distributed system parameters
 - b. Compare to existing system

Conclusion

- Performance increase due to distributed queues and workers?
- CQRS can be implemented in a distributed way
 - Application: PHP and Symfony2
 - Database: MySQL
 - Eventbus: RabbitMQ
 - Eventstore: Elasticsearch
- Lessons learned
 - Making an existing solution distributed is kind of hard (loads of dependencies)
 - Symfony2 is not designed for clustering RabbitMQ
 - Eventual consistency is tricky to test